

A Step-by-Step Guide to Installing OpenStack on CentOS Using the KVM Hypervisor and GlusterFS Distributed File System

Anton Beloglazov Sareh Fotuhi Piraghaj
Mohammed Alrokayan Rajkumar Buyya

Contents

1	Introduction	3
2	Comparison of Open Source Cloud Platforms	3
3	Overview of the OpenStack Cloud Platform	3
4	Existing OpenStack Installation Tools	3
5	Step-by-Step OpenStack Installation	4
5.1	Hardware Setup	4
5.2	Organization of the Installation Package	5
5.3	Configuration Files	5
5.4	Installation Procedure	6
5.4.1	CentOS Installation	6
5.4.2	Network Gateway	9
5.4.3	GlusterFS Distributed Replicated Storage	9
5.4.4	KVM	9
5.4.5	OpenStack	9
5.4.6	Testing of the OpenStack Installation	10
5.5	OpenStack Troubleshooting	10

6 Conclusion	10
7 References	10

1 Introduction

- Cloud Computing [1], [2]
- Public / Private / Hybrid
- Why Open Source Cloud Platforms are Important
- OpenStack / Eucalyptus / CloudStack / OpenNebula
- Complexity of Installing OpenStack
- Our Step-by-Step Scripted Installation Approach
- The purpose is not just having an up and running OpenStack installation, but also learning the steps required to perform the installation from the ground up and understanding the responsibilities and interaction of the OpenStack components.

2 Comparison of Open Source Cloud Platforms

- OpenStack
- Eucalyptus
- CloudStack
- OpenNebula

3 Overview of the OpenStack Cloud Platform

- History
- Features
- Main Services
- Service Interaction

4 Existing OpenStack Installation Tools

- DevStack¹
- Puppet / Chef²
- Difference From our Approach

¹[Http://devstack.org/](http://devstack.org/).

²[Http://docs.openstack.org/trunk/openstack-compute/admin/content/openstack-compute-deployment-tool-with-puppet.html](http://docs.openstack.org/trunk/openstack-compute/admin/content/openstack-compute-deployment-tool-with-puppet.html).

5 Step-by-Step OpenStack Installation

5.1 Hardware Setup

The testbed used for testing the installation scripts consists of the following hardware:

- 1 x Dell Optiplex 745
 - Intel(R) Core(TM)2 CPU (2 cores, 2 threads) 6600 @ 2.40GHz
 - 2GB DDR2-667
 - Seagate Barracuda 80GB, 7200 RPM SATA II (ST3808110AS)
 - Broadcom 5751 NetXtreme Gigabit Controller
- 4 x IBM System x3200 M3
 - Intel(R) Xeon(R) CPU (4 cores, 8 threads), X3460 @ 2.80GHz
 - 4GB DDR3-1333
 - Western Digital 250 GB, 7200 RPM SATA II (WD2502ABYS-23B7A)
 - Dual Gigabit Ethernet (2 x Intel 82574L Ethernet Controller)
- 1 x Netgear ProSafe 16-Port 10/100 Desktop Switch FS116

The Dell Optiplex 745 machine has been chosen to serve as a management host running all the major OpenStack services. The management host is referred to as the *controller* further in the text. The 4 IBM System x3200 M3 servers are used as *compute hosts*, i.e. for hosting VM instances.

Due to specifics of our setup, the only one machine connected to public network and the Internet is one of the IBM System x3200 M3 servers. This server is referred to as the *gateway*. The gateway is connected to the public network via the eth0 network interface.

All the machines form a local network connected through the Netgear FS116 network switch. The compute hosts are connected to the local network via their eth1 network interfaces. The controller is connected to the local network through its eth0 interface. To provide the access to the public network and the Internet, the gateway performs Network Address Translation (NAT) for the hosts from the local network.

5.2 Organization of the Installation Package

The project contains a number of directories, whose organization is explained in this section. The **config** directory includes configuration files, which are used by the installation scripts and should be modified prior to the installation. The **lib** directory contains utility scripts that are shared by the other installation scripts. The **doc** directory comprises the source and compiled versions of the documentation.

The remaining directories directly include the step-by-step installation scripts. The names of these directories have a specific format. The prefix (before the first dash) is the number denoting the order of execution. For example, the scripts from the directory with the prefix *01* must be executed first, followed by the scripts from the directory with the prefix *02*, etc. The middle part of a directory name denotes the purpose of the scripts in this directory. The suffix (after the last dash) specifies the host, on which the scripts from this directory should be executed on. There are 4 possible values of the target host prefix:

- *all* – execute the scripts on all the hosts;
- *compute* – execute the scripts on all the compute hosts;
- *controller* – execute the scripts on the controller;
- *gateway* – execute the scripts on the gateway.

For example, the first directory is named **01-network-gateway**, which means that (1) the scripts from this directory must be executed in the first place; (2) the scripts are supposed to do a network set up; and (3) the scripts must be executed only on the gateway. The name **02-glusterfs-all** means: (1) the scripts from this directory must be executed after the scripts from **01-network-gateway**; (2) the scripts set up GlusterFS; and (3) the scripts must be executed on all the hosts.

The names of the installation scripts themselves follow a similar convention. The prefix denotes the order, in which the scripts should be run, while the remaining part of the name describes the purpose of the script.

5.3 Configuration Files

The **lib** directory contains configuration files used by the installation scripts. These configuration files should be modified prior to running the installation scripts. The configuration files are described below.

configrc: This file contains a number of environmental variables defining various aspects of OpenStack's configuration, such as administration and

service account credentials, as well as access points. The file must be “sourced” to export the variables into the current shell session. The file can be sourced directly by running: `. configrc`, or using the scripts described later. A simple test to check whether the variables have been correctly exported is to `echo` any of the variables. For example, `echo $OS_USERNAME` must output `admin` for the default configuration.

hosts: This file contains a mapping between the IP addresses of the hosts in the local network and their host names. We apply the following host name convention: the compute hosts are named *computeX*, where *X* is replaced by the number of the host. According to the described hardware setup, the default configuration defines 1 **controller** (192.168.0.1), and 4 compute hosts: **compute1** (192.168.0.1), **compute2** (192.168.0.2), **compute3** (192.168.0.3), **compute4** (192.168.0.4). As mentioned above, in our setup one of the compute hosts is connected to the public network and acts as a gateway. We assign to this host the host name **compute1**, and also alias it as **gateway**.

ntp.conf: This file contains a list of Network Time Protocol (NTP) servers to use by all the hosts. It is important to set accessible servers, since time synchronization is important for OpenStack services to interact correctly. By default, this file defines servers used within the University of Melbourne. It is advised to replace the default configuration with a list of preferred servers.

It is important to replace the default configuration defined in the described configuration files, since the default configuration is tailored to the specific setup of our testbed.

5.4 Installation Procedure

5.4.1 CentOS Installation

The installation scripts have been tested with CentOS 6.3³, which has been installed on all the hosts. The CentOS installation mainly follows the standard process described in detail in the Red Hat Enterprise Linux 6 Installation Guide [3]. The steps of the installation process that differ from the standard are discussed in this section.

Network Configuration. The simplest way to configure network is during the OS installation process. As mentioned above, in our setup, the gateway is connected to two networks: to the public network through the `eth0` interface; and to the local network through the `eth1` interface. Since in our setup the public

³[Http://www.centos.org/](http://www.centos.org/).

network configuration can be obtain from a DHCP server, in the configuration of the eth0 interface it is only required to enable automatic connection by enabling the “Connect Automatically” option. We use static configuration for the local network; therefore, eth1 has be configured manually. Apart from enabling the “Connect Automatically” option, it is necessary to configure IPv4 by adding an IP address and netmask. According to the configuration defined in the `hosts` file described above, we assign 192.168.0.1/24 to the gateway.

One difference in the network configuration of the other compute hosts (`compute2`, `compute3`, and `compute4`) from the gateway is that eth0 should be kept disabled, as it is unused. The eth1 interface should be enabled by turning on the “Connect Automatically” option. The IP address and netmask for eth1 should be set to 192.168.0.*X*/24, where *X* is replaced by the compute host number. The gateway for the compute hosts should be set to 192.168.0.1, which the IP address of the gateway host. The controller is configured similarly to the compute hosts with the only difference that the configuration should be done for eth0 instead of eth1, since the controller has only one network interface.

Hard Drive Partitioning. The hard drive partitioning scheme is the same for all the compute hosts, but differs for the controller. Table 1 shows the partitioning scheme for the compute hosts. `vg_base` is a volume group comprising the standard operating system partitions: `lv_root`, `lv_home` and `lv_swap`. `vg_gluster` is a special volume group containing a single `lv_gluster` partition, which is dedicated to serve as a GlusterFS brick.

Device	Size (MB)	Mount Point / Volume	Type
<i>LVM Volume Groups</i>			
vg_base	20996		
lv_root	10000	/	ext4
lv_swap	6000		swap
lv_home	4996	/home	ext4
vg_gluster	216972		
lv_gluster	216972	/export/gluster	xfs
<i>Hard Drives</i>			
sda			
sda1	500	/boot	ext4
sda2	21000	vg_base	PV (LVM)
sda3	216974	vg_gluster	PV (LVM)

Table 1: Partitioning scheme for the compute hosts

Table 2 shows the partitioning scheme for the controller. It does not include a **vg_gluster** volume group since the controller is not going to be a part of the GlusterFS volume. However,

Device	Size (MB)	Mount Point / Volume	Type
<i>LVM Volume Groups</i>			
nova-volumes	29996		
Free	29996		
vg_base	16996		
lv_root	10000	/	ext4
lv_swap	2000		swap
lv_home	4996	/home	ext4
vg_images	28788		
lv_images	28788	/var/lib/glance/images	ext4
<i>Hard Drives</i>			
sda			
sda1	500	/boot	ext4
sda2	17000	vg_base	PV (LVM)
sda3	30000	nova-volumes	PV (LVM)
sda4	28792		Extended
sda5	28788	vg_images	PV (LVM)

Table 2: Partitioning scheme for the controller

5.4.2 Network Gateway

5.4.3 GlusterFS Distributed Replicated Storage

All nodes

Controller

All nodes

5.4.4 KVM

5.4.5 OpenStack

All nodes

Controller

Compute nodes

Network Gateway

Controller

5.4.6 Testing of the OpenStack Installation

5.5 OpenStack Troubleshooting

6 Conclusion

7 References

- [1] M. Armbrust, A. Fox, R. Griffith, A. D. Joseph, R. Katz, A. Konwinski, G. Lee, D. Patterson, A. Rabkin, I. Stoica, and others, “A view of cloud computing,” *Communications of the ACM*, vol. 53, pp. 50–58, 2010.
- [2] R. Buyya, C. S. Yeo, S. Venugopal, J. Broberg, and I. Brandic, “Cloud computing and emerging IT platforms: Vision, hype, and reality for delivering computing as the 5th utility,” *Future Generation computer systems*, vol. 25, pp. 599–616, 2009.
- [3] R. Landmann, J. Reed, D. Cantrell, H. D. Goede, and J. Masters, “Red Hat Enterprise Linux 6 Installation Guide,” 2012.